

# OS Lab No. 1

Date: 04 / 08 / 2025

## Team Members:

1. Kashyap Pansuriya - 230123079
2. Gopal Singh - 230123083
3. Aditya Arun Gupta - 230123078
4. Nilarnab Sutradhar - 230123041

## Task 1: Prime PID Assign

Implement the feature to assign prime number PIDs to newly created processes in xv6-riscv.

### Snippet 1: kernel/proc.c

```
// Allocate only prime pid to processes
static int
isprime(int n)
{
    if(n < 2) return 0;
    if(n % 2 == 0) return n == 2;
    for(int d = 3; d*d <= n; d += 2)
        if(n % d == 0)
            return 0;
    return 1;
}

static int
nextprime(int start)
{
    int p = start;
    while(!isprime(p)) {
        p++;
    }
    return p;
}

int
allocpid()
{
    int pid;

    acquire(&pid_lock);
    pid = nextprime(nextpid);
    nextpid = pid + 1;
    release(&pid_lock);

    return pid;
}
```

- isprime() and nextprime() to determine prime numbers for PID allocation.
- Modified allocpid() so it skips non-prime PIDs.

## Snippet 2: user/primepidTest.c

```
#include "kernel/types.h"
#include "user/user.h"

#define N 10000

int
main(void)
{
    for(int i = 0; i < N; i++){
        int pid = fork();
        if(pid < 0){
            printf("fork failed\n");
            exit(1);
        }
        if(pid == 0){
            printf("child %d: my PID %d\n", i+1, getpid());
            exit(0);
        }
        else {
            wait(0);
        }
    }

    exit(0);
}
```

- Test to verify newly spawned processes receive prime number PIDs.
- Test can be invoked from xv6 terminal using **primepidTest** command.

## Task 2: Implement top Command

Expose process information via a new syscall and build a user-space `top` command.

### Snippet 1: kernel/proc.h

```
#include "procinfo.h"
...

int getprocinfo(struct procinfo *info);
...
```

### Snippet 2: kernel/proc.c

```
int getprocinfo(struct procinfo *info)
{
    struct proc *p;
    int idx = 0;

    for (int i = 0; i < NPROC; i++)
    {
        p = &proc[i];
        acquire(&p->lock);

        if (p->state != UNUSED)
        {
            info[idx].pid = p->pid;
            info[idx].state = p->state;
            info[idx].ticks = p->ticks;

            safestrcpy(info[idx].name, p->name, sizeof(info[idx].name));
            idx++;
        }

        release(&p->lock);
    }

    return idx;
}
```

- Defines helper function to get process information.
- **getprocinfo** function fills up struct `procinfo` which stores pid, state, ticks and name of the process

### Snippet 3: kernel/trap.c

```
...
struct spinlock tickslock;
uint ticks;
...

void
trapinit(void)
{
    initlock(&tickslock, "time");
}

void
clockintr()
{
    if(cpuid() == 0){
        acquire(&tickslock);
        ticks++;
        wakeup(&ticks);
        release(&tickslock);
    }

    w_stimecmp(r_time() + 1000000);
}

int
devintr()
{
    ...
    if(myproc() != 0 && myproc()->state == RUNNING) myproc()->ticks++;
    ...
}
```

- Set up a ‘ticks’ counter protected by **tickslock** in **trapinit()** so we can safely track timer ticks.
- In **clockintr()**, we grab the lock, bump the global ‘ticks’, wake any waiting processes, and release the lock, then schedule the next interrupt.
- After each timer interrupt, if a user process is running, we do **myproc()-ticks++** in **usertrap()** to count one CPU-time tick for that process.
- Together, these changes let us keep both a system-wide tick count and individual process tick counts for the **top** command

## Snippet 4: kernel/procinfo.h

```
#pragma once
#include "kernel/types.h"

#ifndef STATE_DEFINE
#define STATE_DEFINE
enum procstate { UNUSED, USED, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
#endif

#ifndef PROCINFO
#define PROCINFO
struct procinfo
{
    int pid;
    enum procstate state;
    uint ticks;
    char name[16];
};
#endif
```

## Snippet 5: kernel/sysproc.c

```
...
extern struct proc proc[NPROC];
...

uint64
sys_getprocinfo(void)
{
    uint64 addr;
    struct procinfo info[NPROC];
    int count;

    if (argaddr(0, &addr) < 0)
        return -1;

    count = getprocinfo(info);

    if (copyout(myproc()->pagetable, addr, (char *)info, sizeof(info)) < 0)
        return -1;

    return count;
}
```

- Creating a interface to be able to use getprocsinfo as a system call from user-space

## Snippet 6: kernel/syscall.c and kernel/syscall.h

```
...
extern uint64 sys_getprocinfo(void);
...

static uint64 (*syscalls[])(void) = {
...
[SYS_getprocinfo] sys_getprocinfo,
...
};
```

## Snippet 7: user/user.h

```
...
int getprocinfo(void *);
...
```

## Snippet 8: user/usys.pl

```
...
entry("getprocinfo");
...
```

- Some necessary changes in order to define a new custom system call

## Snippet 9: kernel/top.c

```
#include "kernel/types.h"
#include "kernel/stat.h"
#include "user.h"
#include "kernel/procinfo.h"

#define NPROC 64

void print_procs(){
    struct procinfo arr[NPROC];
    int n = getprocinfo(arr);

    if (n < 0)
    {
        printf("Failed to retrieve process info\n");
        exit(1);
    }

    printf("PID \t STATE \t\t TICKS \t NAME\n");
    printf("-----\n");

    for (int i = 0; i < n; i++)
    {
        char *state;

        switch (arr[i].state)
        {
            case UNUSED: continue; // Skip UNUSED
            case USED: state = "USED"; break;
            case SLEEPING: state = "SLEEPING"; break;
            case RUNNABLE: state = "RUNNABLE"; break;
            case RUNNING: state = "RUNNING"; break;
            case ZOMBIE: state = "ZOMBIE"; break;
            default: continue; // Skip unknown
        }

        printf("%d \t %s \t %d \t %s\n", arr[i].pid, state, arr[i].ticks, arr[i].name);
    }

    printf("\n");
}
```

- Print function to print the formatted table of processes as per requirement



## Conclusion and References

We implemented the Prime PID assignment and a user-level `top` utility in xv6-riscv, adding kernel data structures, syscalls, and user-space support. The tests demonstrate correct PID assignment, tick accounting, and process-table rendering.

### References:

- xv6-riscv source code (MIT)
- Lab assignment description